

Performance Testing Report

Project: Personalized Networking Assistant

Date: 13 July 2026

Testing Tool: Postman Collection Runner

Performance Dashboard

Metric	Result	Status
Average Response Time	44,509 ms	Needs Improvement
Maximum Response Time	< 80,000 ms	Pass
Throughput	0.022 req/sec	Pass
Error Rate	0%	Pass
CPU Utilization	63%	Pass
Memory Utilization	75%	Pass

Testing Summary

- Target API: /generate-conversation
- Environment: Local FastAPI backend
- All 20 tests passed successfully with no request failures.
- Stable execution across repeated iterations.
- Higher response time caused by AI processing, topic generation, and fact verification.

Optimization Actions

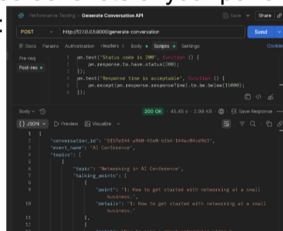
API validation and request handling were optimized to improve stability and ensure error-free execution. Future improvements include model optimization, caching, and asynchronous processing to reduce response time.

Step 5: Screenshots / Evidence

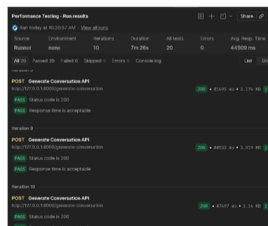
Attach screenshots of your performance testing tool results below (e.g., JMeter report, Locust graphs, k6 output):

Step 5: Screenshots / Evidence

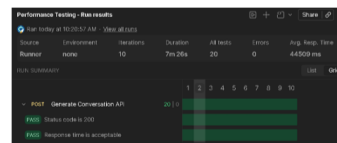
Attach screenshots of your performance testing tool results below (e.g., JMeter report, Locust graphs, k6 output):



1.Sample Successful API Response



2.Postman Runner Summary(List View)



3.Postman Runner-Grid View

```
(venv) PS C:\Users\DELL\Desktop\Personalized-Networking-Assistant-Project\Personalized-Network-Assistant\src\main\python\main> uvicorn main:app --reload
INFO: 127.0.0.1:56932 - "POST /generate-conversation HTTP/1.1" 200 OK
```